

Airline Data Streaming

Intro

Data streaming has become an increasingly important technique in recent years due to both size and the dynamic nature of data. In this assignment, we work with data in “blocks”. In other words, we don’t read the entire data into memory but rather read a subset of the data, process it immediately, discard it right away and continue with the next subsets of the data. Therefore, we can have very large scale of data streaming into our memory and extract statistical information at the same time.

We will compare the speed of processing data using Unix shell tools, pure R tools and combination of both tools. We found that a connection between R and Unix is an efficient way to stream and process the data. Database tools such as SQLite is also explored to improve the efficiency of processing large scale of data. We will compute the number of flights leaving each of the airports (LAX, OAK, SFO, SMF) from 1987 to 2008. We will also compute the mean and standard deviation of the arrival delay times for flights departing from these four airports.

Data Structure

The data sets are constructed by 22 csv files (totally 12 Gigabytes), which contain domestic airline flights summaries as well as arrival and departure performance each year from 1987 to 2008. We have 29 columns in each data set for different years. We have particular interests in the 15th column for arrival delay (ARR_DELAY) and 17th column for departure airport (ORIGIN). Other relevant time variable includes YEAR, MONTH, DAY_OF_MONTH, DAY_OF_WEEK.

The data are available at <http://eeyore.ucdavis.edu/stat242/data/>. Please see Appendix I for R codes to download and uncompressed files programmatically.

Tools Exploration

The tools used in this assignment include Unix shell tools, R tools and database engine. We compare the speed of counting flights leaving each of the airports, specifically LAX, OAK, SFO and SMF for year 2008 to find the most efficient way to obtain statistical information (e.g. counts, means and standard deviations) for the entire data sets. Following is the system time needed to for counting:

	Shell Scripting		R Streaming	
	i. Multi passes (wc -l)	ii. One pass (sort uniq -c)	i. readLines (file(csv))	ii. readLines (pipe(egrep))
user	67.933	62.077	268.294	19.476
system	0.833	0.242	1.229	0.061
elapsed	55.612	61.230	269.525	57.478

Table -1. System.time for counts of flight leaving LAX, OAK, SFO and SMF in 2008

We found that Unix shell tools are extremely fast in taking subsets of data using “cut” and “grep” commands, which select the specific column and search with regular expression. “sort” and “uniq -c” is easier to write without doing multiple passes but it takes longer to sort given the data set is larger enough.

```
system.time(SHcounts.wc<-system("for airport in LAX OAK SFO SMF; do cut -f 17
-d , 2008.csv | grep $airport | wc -l; done",intern=TRUE))

system.time(SHcounts.uniq<-system("egrep '([0-9]|NA),(LAX|OAK|SFO|SMF),[A-
Z]' 2008.csv | cut -f 17 -d , | sort | uniq -c",intern=TRUE))
```

Reading data in blocks directly from csv file in R is the slowest way to stream and process the data. However, prepare subset data in the Unix environment with regular expression significantly reduce the processing time from 269.525 seconds to 57.478 seconds. This is close enough to counting lines in shell scripting. Codes are available in Appendix II.

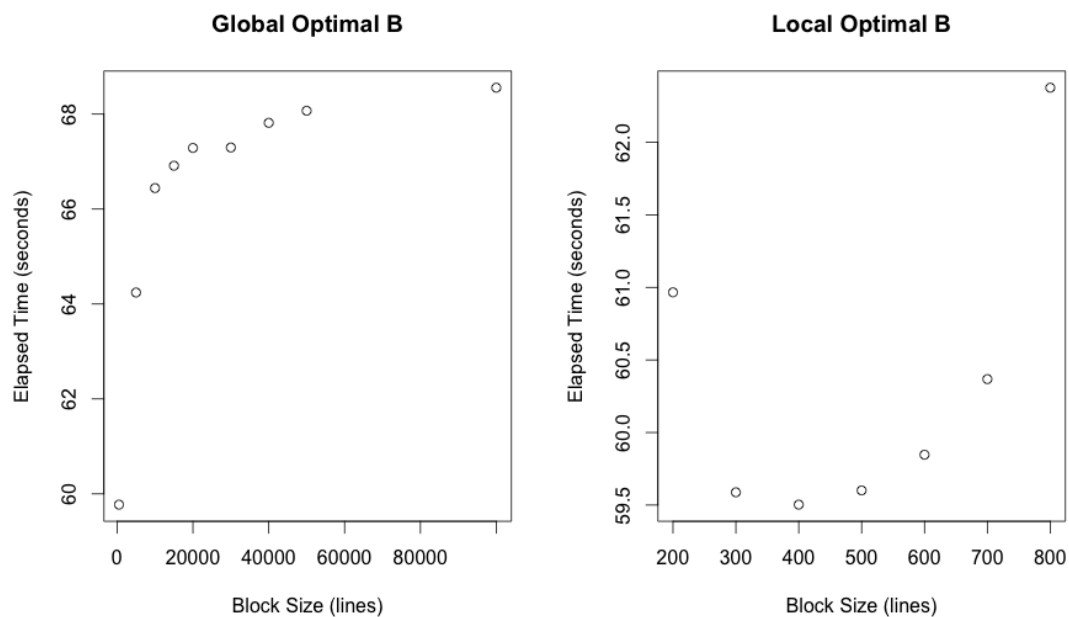
```
con<-pipe("egrep '([0-9]|NA),(LAX|OAK|SFO|SMF),[A-Z]' 2008.csv","r")

rcount<-function(B,con,year){
  B=as.integer(B)
  # set up counters
  Rcounts<-structure(c(data=rep(0,4)),names=c("LAX","OAK","SFO","SMF"))
  # reading blocks
  while(TRUE){
    txt=readLines(con,n=B)
    if (length(txt)==0)
      break
    # counting
    temp=sapply(strsplit(txt,","),"[",17)
    update<-as.numeric(table(temp)[names(Rcounts)])
    update[is.na(update)]<-0
    Rcounts<-Rcounts+update
  }
  Rcounts
}

system.time(RSHcount<-rcount(400L,con,"2008"))
```

We are interested in exploring the best block size that is highly relevant to the performance for reading data while Unix is processing data line by line. The parameterized function and plot function in R are handy for us to find the optimal block size for streaming data from Unix to R. The optimal block size for reading lines directly from csv files and connection from Unix are different, specifically 12000L versus 4000L, because R is slower than Unix shell tools in taking the subsets of data and thus would be better of in dealing with larger block size in counting. But if we stream data indirectly from the connection of Unix using pipe(), we are more efficient in counting targeted observations with smaller block size than that in reading directly from csv files.

Following are plots for the exploration of optimal block size in the case of streaming prepared data from Unix to R :



The optimal B is 400 lines for an elapsed time less than 60 seconds if we pre-process data in Unix and then stream to R. We can expect to improve the efficiency by interfacing to C code because we could do faster while loops.

With the guarantee of speed, we would like to use R for its good extensibility with handy statistical tools. Functions such as `table()`, `sum()` and `var()` will help us to obtain the counts, mean and standard deviation.

Statistics: counts, means and standard deviation

Grounded on the data streaming in R reading from Unix, a function `rstats()` is developed to get the counts, obtain the sum and cumulate the variance for airlines leaving LAX, OAK, SFO and SMF. The mean is the sum divided by counts while the standard deviation is the square root of variance divided by counts.

(I) Outliers

Notice that the counts here are different from the counts obtained in the previous section (Tools Exploration) because we exclude the observations with NA on arrival delay. We will obtain more accurate mean and standard deviation for the arrival delay time by excluding the NA observations.

The counts of ORIGIN for each airport in 2008:

```
> LAX 184048; > OAK 53679; > SFO 118635; > SMF 45364; #NA included
> LAX 181308; > OAK 52818; > SFO 116029; > SMF 44907; #NA excluded
```

The counts of ORIGIN for each airport from 1987 to 2008:

```
> LAX 3947365; > OAK 1134520; > SFO 2605579; > SMF 790378; #NA included
> LAX 3879885; > OAK 1121028; > SFO 2551563; > SMF 782171; #NA excluded
```


(II) AWK Validation

We use awk with under UNIX to do a quick validation of the sum of arrival delay time. It would print the counts and cumulated arrival delay time. The statistical results in R are validated by this AWK validation for 1987, 2001 and 2008.

```
system.time(cmean<-system("egrep '([0-9]|NA),LAX,[A-Z]' 2008.csv | cut -f 15 -d , | awk 'BEGIN {s=0; c=0}; {s=s+$1;c=c+1}; END {print c,s, s/c}'",intern=TRUE))
```

(III) Results

The statistical results are constructed as arrays. Columns are the statistics including counts, sum, variance; rows are the airports of our interests, specifically for LAX, OAK, SFO and SMF; the third dimension is the year from 1987 to 2008. Please see Appendix III for the statistics obtained from data streaming.



We can use apply function to obtain the mean and std.dev for each year. Results are available in Appendix IV.

```
#counts
apply(RS.ALL.ARRAY,c(3),function(x){x[,1]})
#means
apply(RS.ALL.ARRAY,c(3),function(x){x[,2]/x[,1]})
#standard deviation
apply(RS.ALL.ARRAY,c(3),function(x){sqrt(x[,3])/x[,1]})
```

To obtain the statistics for the entire data set from 1987 to 2008, we can use apply function to sum up the counts, sum and variance through out the time.

```
TOTAL<-apply(RS.ALL.ARRAY,c(1,2),sum)
TOTAL[,1]#counts
TOTAL[,2]/TOTAL[,1] #mean
sqrt(TOTAL[,3]/TOTAL[,1]) #std.dev
```

The final statistical results for the entire data set is shown in the table below:

	Counts	Means	Std.Dev
LAX	3879885	6.011921	26.31361
OAK	1121028	5.027054	21.09526
SFO	255156	8.008065	29.06354
SMF	782171	5.342205	24.14323

Conclusion

In this assignment, we work with data in “blocks” to break through the limitation of memory and improve speed of data processing. Learning data streaming is an important step to learn parallel computing which carry out the calculations simultaneously. If data streaming is processing data in an time series way, parallel computing is processing data from an cross sectional prospective.

Besides the idea of reading data in blocks, I learn how to use Unix shell tools to pre-process data and open connections via `pipe()` to stream the data into R. Basic commands such as `grep`, `cut`, `uniq`, `sort` and the redirection operators in the shell are very handy in pre-processing data. In R, I learn how to open connections include `pipe`, `file`, `url`, `bzfile`, `gzfile`, `xzfile`.

Character encoding is an issue while dealing with the csv file for 2001 and 2002. It is important to read the data with correct encoding which pairs each character from a given repertoire particular bit patterns.

URL for the repository on GitHub:
<https://github.com/ykangxie/airport.git>

Appendix I.

```
#####  
# Download and Unzip Files  
#####  
  
#/*get fileURI*/  
#txt = readLines("http://eeyore.ucdavis.edu/stat242/data")  
#ll = grep("csv.bz2", txt) #line.number containing "csv.bz2"  
#files = gsub(".*([0-9]{4}.(csv).bz2).*", "\\1", txt[ll]) #file names  
#fileUrl<-sprintf("http://eeyore.ucdavis.edu/stat242/data/%s", files) #fileUrl  
  
#/*download bz2 files*/  
#destFile<-sprintf("./Desktop/airport/%s",files)  
#mapply(function(url,output){download.file(url,destfile=output,method="curl")  
,fileUrl,destFile)  
  
#/*unzip bz2 files*/  
#cmd<-sprintf("cd ./Desktop/airport; bunzip2 -d %s", files)  
#sapply(cmd,system)
```

Appendix II. Tools Exploration

```
files = gsub(".*([0-9]{4}).csv.*", "\\1", list.files()[grep(".*([0-9]{4}).csv.*",
list.files())])
year = gsub(".*([0-9]{4}).csv.*", "\\1", files[grep(".*([0-9]{4}).csv.*", files)])
```

```
#####
# Tools Exploration- system.time(counts) for 2008.csv
#####
setwd("./Desktop/airport")
```

```
dept=c("LAX","OAK","SFO","SMF")
year="2008"
```

#(I) Shell Scripting

#1.1 /grep & wc -l/ [multiple passes]

```
system.time(SHcounts.wc<-system("for airport in LAX OAK SFO SMF; do cut -f 17
-d , 2008.csv | grep $airport | wc -l; done",intern=TRUE))
#user system elapsed
#67.933 0.833 55.612
```

#1.2 /egrep & (sort + uniq -c)/ [one pass]

```
system.time(SHcounts.uniq<-system("egrep '([0-9]|NA),(LAX|OAK|SFO|SMF),[A-
Z]' 2008.csv | cut -f 17 -d , | sort | uniq -c",intern=TRUE))
#user system elapsed
#62.077 0.242 61.230
```

#(II) R Streaming (csv files)

```
rstreaming<-function(B,csvfile="2008.csv"){
  B=as.integer(B)
  con=file(csvfile,"r")
  Rcounts<-structure(c(data=rep(0,4)),names=c("LAX","OAK","SFO","SMF"))
  while(TRUE){
    txt=readLines(con,n=B)
    if (length(txt)==0)
      break
    temp=sapply(strsplit(txt,""),"[[",17)
    update<-as.numeric(table(temp)[names(Rcounts)])
    update[is.na(update)]<-0
    Rcounts<-Rcounts+update
  }
  Rcounts
}
```

```
system.time(rstreaming(400L,"2008.csv")) #268.294 1.229 269.525
system.time(rstreaming(12000L,"2008.csv")) #255.774 0.816 256.571
```

#(III) R Streaming from Unix via pipe()

```
con<-pipe("egrep '([0-9]|NA),(LAX|OAK|SFO|SMF),[A-Z]' 2008.csv","r")
rcount<-function(B,con,year){
  B=as.integer(B)
  Rcounts<-structure(c(data=rep(0,4)),names=c("LAX","OAK","SFO","SMF"))
  while(TRUE){
    txt=readLines(con,n=B)
    if (length(txt)==0)
      break
    temp=sapply(strsplit(txt","),"[",17)
    update<-as.numeric(table(temp)[names(Rcounts)])
    update[is.na(update)]<-0
    Rcounts<-Rcounts+update
  }
  Rcounts
}
system.time(RSHcount<-rcount(400L,con,"2008"))
#user system elapsed
#19.476 0.061 57.478
```


Appendix III. Function rstats()

```

#rstats
rstats<-function(B,con,year){
  B=as.integer(B)
  #set up counters
  Rcounts<-structure(c(data=rep(0,4)),names=c("LAX","OAK","SFO","SMF"))
  Rsum<-structure(c(data=rep(0,4)),names=c("LAX","OAK","SFO","SMF"))
  Rsumvar<-structure(c(data=rep(0,4)),names=c("LAX","OAK","SFO","SMF"))
  #reading blocks
  while(TRUE){
    txt=readLines(con,n=B)
    if (length(txt)==0)
      break
    #ArrDelay & Airport
    ArrDelay=as.numeric(sapply(strsplit(txt,""),"[[",15))
    Origin=as.factor(sapply(strsplit(txt,""),"[[",17))
    #dropNA
    dropNA<-!is.na(ArrDelay)
    ArrDelay<-ArrDelay[dropNA]
    Origin<-Origin[dropNA]
    #1.1 Rcounts
    update<-as.numeric(table(temp)[names(Rcounts)])
    #update<-tapply(ArrDelay,Origin,length)[names(Rcounts)] #table() is faster than tapply(length)
    update[is.na(update)]<-0
    Rcounts<-Rcounts+as.numeric(update)
    #1.2 Rsum
    update<-tapply(ArrDelay,Origin,sum)[names(Rsum)]
    update[is.na(update)]<-0
    Rsum<-Rsum+as.numeric(update)
    #1.3 Rvar
    update<-tapply(ArrDelay,Origin,function(x){var(x)*length(x)})[names(Rsumvar)]
    update[is.na(update)]<-0
    Rsumvar<-Rsumvar+as.numeric(update)
  }
  cbind(Rcounts,Rsum,Rsumvar)
}

```

Appendix IV. Codes for All-Years Statistics

```
#####
# All Years: 1987-2012.csv
#####
source('rstats.R')
Sys.setlocale(locale="C")

# i. Each Year (array)
year = gsub(".*([0-9]{4}).csv.*", "\\1", files[grep(".*([0-9]{4}).csv.*", files)])
cmds<-sprintf("egrep '([0-9]|NA),(LAX|OAK|SFO|SMF),[A-Z]' %s.csv",year)
system.time(RS.ALL<-lapply(cmds,function(cmd){con<-pipe(cmd,"r")
                        tmp<-rstats(400L,con)
                        close(con)
                        tmp}))
#user system elapsed #1824.276 10.011 1285.353
names(RS.ALL)<-year
RS.ALL.ARRAY<-array(unlist(RS.ALL), dim = c(nrow(RS.ALL[[1]]),
ncol(RS.ALL[[1]]), length(RS.ALL))) #reconstruction: (4 x 3 x 22)
rownames(RS.ALL.ARRAY)<-rownames(RS.ALL[[1]])
colnames(RS.ALL.ARRAY)<-colnames(RS.ALL[[1]])
dimnames(RS.ALL.ARRAY)[[3]]<-year
#counts
apply(RS.ALL.ARRAY,c(3),function(x){x[,1]})
#means
apply(RS.ALL.ARRAY,c(3),function(x){x[,2]/x[,1]})
#standard deviation
apply(RS.ALL.ARRAY,c(3),function(x){sqrt(x[,3])/x[,1]})

# ii. All Year (matrix)
TOTAL<-apply(RS.ALL.ARRAY,c(1,2),sum)
TOTAL[,1]#counts
TOTAL[,2]/TOTAL[,1] #mean
sqrt(TOTAL[,3]/TOTAL[,1]) #std.dev
TOTAL.STATS<-
structure(data.frame(cbind(TOTAL[,1],TOTAL[,2]/TOTAL[,1],sqrt(TOTAL[,3]/T
OTAL[,1]))),names=c("counts", "means", "std.dev"))

# iii. All Years (awk-validation)
# notes: the awk validation is a quick validation which has a larger n because of
the inclusion of NA ArrDelay observations

system.time(cmean2<-system("egrep '([0-9]|NA),LAX,[A-Z]' [12]*.csv | cut -f 15 -
d , | awk 'BEGIN {s=0; c=0}; {s=s+$1;c=c+1}; END {print c,s,s/c}'",intern=TRUE))

# counts, sum & means for (LAX [12]*.csv)
#[1] "3947365 5.90915"
```

Appendix V. Statistical Information for Each Year

> RS.ALL.ARRAY

SMF 29709 69666 8026645

,, 1987

,, 1993

Rcounts Rsum Rsumvar
LAX 44718 698222 46629543
OAK 8256 72969 3989204
SFO 34387 576020 33625962
SMF 6797 58956 4028147

Rcounts Rsum Rsumvar
LAX 151020 445402 58514989
OAK 39316 82182 7731066
SFO 116523 411976 49284859
SMF 29460 89367 8990148

,, 1988

,, 1994

Rcounts Rsum Rsumvar
LAX 168315 494824 72032151
OAK 28069 119662 8219810
SFO 130309 987464 66265631
SMF 27202 77202 8481110

Rcounts Rsum Rsumvar
LAX 151919 706131 56519438
OAK 44138 182071 9094462
SFO 117053 528507 48494829
SMF 28717 123118 8317834

,, 1989

,, 1995

Rcounts Rsum Rsumvar
LAX 161229 1026141 82462028
OAK 27849 170309 9736698
SFO 124400 909651 75711416
SMF 25118 136822 11395138

Rcounts Rsum Rsumvar
LAX 175894 1674907 117273324
OAK 64524 364986 22515629
SFO 126283 1128107 86908794
SMF 35628 176701 15460096

,, 1990

,, 1996

Rcounts Rsum Rsumvar
LAX 168100 1017515 77791138
OAK 36207 125261 10410748
SFO 128921 820036 61602123
SMF 24919 107973 8888563

Rcounts Rsum Rsumvar
LAX 179193 1840182 144844249
OAK 59669 346971 22738325
SFO 132074 1639385 138390756
SMF 36424 230301 22428645

,, 1991

,, 1997

Rcounts Rsum Rsumvar
LAX 155002 1059089 80226036
OAK 39793 159760 11079901
SFO 122423 1098101 66316359
SMF 26805 117460 8886301

Rcounts Rsum Rsumvar
LAX 183447 1496839 131055064
OAK 57718 302076 19093393
SFO 135199 1371023 106307297
SMF 36767 229068 18147255

,, 1992

,, 1998

Rcounts Rsum Rsumvar
LAX 153756 668930 61829138
OAK 36497 55264 6543787
SFO 121240 550297 44178952

Rcounts Rsum Rsumvar
LAX 178421 1145846 153113710
OAK 55057 364247 27315707

STA 242 HWK4 Yingkang Xie

SFO 133557 1748564 181440746
SMF 36073 252998 21311474

,,1999

Rcounts Rsum Rsumvar
LAX 185812 1612380 156361588
OAK 54725 366694 22596092
SFO 131444 1206069 153878008
SMF 36308 261399 20734114

,,2000

Rcounts Rsum Rsumvar
LAX 203598 2430501 237508835
OAK 56106 616442 40836851
SFO 127532 1857824 200120593
SMF 38301 421899 31972521

,,2001

Rcounts Rsum Rsumvar
LAX 124679 649956 82336267
OAK 47206 319414 22143188
SFO 26523 80109 22390373
SMF 26216 148242 13099860

,,2002

Rcounts Rsum Rsumvar
LAX 170178 44932 102203954
OAK 59085 278342 27347972
SFO 86901 40047 61185023
SMF 37570 140758 20426443

,,2003

Rcounts Rsum Rsumvar
LAX 221122 169021 126010691
OAK 66662 99080 29496751
SFO 120866 93931 70251106
SMF 46338 114649 23423413

,,2004

Rcounts Rsum Rsumvar
LAX 229731 980717 154611681
OAK 70049 320588 38459632
SFO 127990 526323 88707721

SMF 48110 328413 37680364

,,2005

Rcounts Rsum Rsumvar
LAX 228122 1081568 153685230
OAK 69652 377980 38562684
SFO 127055 813418 118866372
SMF 50161 314357 35258171

,,2006

Rcounts Rsum Rsumvar
LAX 230381 1341989 195359125
OAK 73761 383264 49994165
SFO 129245 1234712 143417625
SMF 53218 265041 41288074

,,2007

Rcounts Rsum Rsumvar
LAX 233940 1779392 221358053
OAK 73871 419482 43786892
SFO 135609 1582832 172304007
SMF 57423 364578 51483842

,,2008

Rcounts Rsum Rsumvar
LAX 181308 961077 174730165
OAK 52818 108424 27175568
SFO 116029 1228686 165629128
SMF 44907 149550 36196022

Appendix VI. Statistical Results for Each Year

```
> #counts
> apply(RS.ALL.ARRAY,c(3),function(x){x[,1]})
      1987  1988  1989  1990  1991  1992  1993  1994  1995  1996  1997  1998  1999
LAX 44718 168315 161229 168100 155002 153756 151020 151919 175894 179193 183447 178421 185812
OAK  8256  28069  27849  36207  39793  36497  39316  44138  64524  59669  57718  55057  54725
SFO 34387 130309 124400 128921 122423 121240 116523 117053 126283 132074 135199 133557 131444
SMF  6797  27202  25118  24919  26805  29709  29460  28717  35628  36424  36767  36073  36308
      2000  2001  2002  2003  2004  2005  2006  2007  2008
LAX 203598 124679 170178 221122 229731 228122 230381 233940 181308
OAK  56106  47206  59085  66662  70049  69652  73761  73871  52818
SFO 127532 26523  86901 120866 127990 127055 129245 135609 116029
SMF  38301  26216  37570  46338  48110  50161  53218  57423  44907

> #means
> apply(RS.ALL.ARRAY,c(3),function(x){x[,2]/x[,1]})
      1987  1988  1989  1990  1991  1992  1993  1994  1995  1996
LAX 15.613891 2.939869 6.364494 6.053034 6.832744 4.350594 2.949291 4.648076 9.522252 10.269274
OAK  8.838299 4.263137 6.115444 3.459580 4.014776 1.514207 2.090294 4.125040 5.656593  5.814929
SFO 16.751098 7.577865 7.312307 6.360764 8.969728 4.538906 3.535577 4.515109 8.933166 12.412625
SMF  8.673827 2.838100 5.447169 4.332959 4.382018 2.344946 3.033503 4.287286 4.959610  6.322782
      1997  1998  1999  2000  2001  2002  2003  2004  2005
LAX  8.159517  6.422148  8.677480 11.93774  5.213035 0.2640294 0.7643789 4.268980 4.741182
OAK  5.233653  6.615816  6.700667 10.98710  6.766386 4.7108742 1.4863040 4.576625 5.426693
SFO 10.140778 13.092268 9.175535 14.56751  3.020360 0.4608347 0.7771499 4.112220 6.402094
SMF  6.230261  7.013500  7.199488 11.01535  5.654638  3.7465531 2.4741896 6.826294 6.266960
      2006  2007  2008
LAX 5.825085  7.606190  5.300798
OAK 5.196025  5.678575  2.052785
SFO 9.553267 11.672028 10.589473
SMF 4.980289  6.348989  3.330216

> #standard deviation
> apply(RS.ALL.ARRAY,c(3),function(x){sqrt(x[,3])/x[,1]})
      1987  1988  1989  1990  1991  1992  1993  1994
LAX 0.1527032 0.05042436 0.05632275 0.05246834 0.05778570 0.05114045 0.05065229 0.04948651
OAK 0.2419209 0.10214190 0.11204597 0.08911451 0.08364910 0.07009022 0.07072132 0.06832440
SFO 0.1686332 0.06246975 0.06994558 0.06087993 0.06651925 0.05482287 0.06024834 0.05949290
SMF 0.2952809 0.10705957 0.13439241 0.11964242 0.11121022 0.09536280 0.10177724 0.10043060
      1995  1996  1997  1998  1999  2000  2001  2002
LAX 0.06156709 0.06716292 0.06240455 0.06935233 0.06729631 0.07569486 0.07277836 0.05940602
OAK 0.07353952 0.07991540 0.07570600 0.09492786 0.08686221 0.11389819 0.09968341 0.08850857
SFO 0.07382220 0.08907098 0.07626198 0.10085577 0.09437295 0.11092431 0.17840541 0.09001146
SMF 0.11036075 0.13002111 0.11586367 0.12797481 0.12541238 0.14763126 0.13805969 0.12029711
      2003  2004  2005  2006  2007  2008
LAX 0.05076586 0.05412545 0.05434366 0.06066947 0.06359796 0.07290662
OAK 0.08147207 0.08853207 0.08915589 0.09585899 0.08957732 0.09869777
SFO 0.06934616 0.07358761 0.08580997 0.09265898 0.09679639 0.11091794
SMF 0.10444504 0.12759163 0.11837601 0.12074072 0.12495384 0.13397272
```